


NEURAL NETWORKS HOMEWORK 2

Convolutional Neural Networks for MNIST and CIFAR data



ENPM 808Y University of Maryland College Park
Timothy Werder

Table of Contents

Problem 1: MNIST	2
Confusion Table	2
Accuracy Table	2
Hyper Parameters	3
Training and Validation loss	3
Discussion.....	4
Problem 2: CIFAR	5
Confusion Table	5
Accuracy Table	5
Hyper Parameters	5
Training and Validation loss	6
Discussion.....	7
Key takeaways.....	8

Problem 1: MNIST

Confusion Table

Train Label	Train Predicted	Test Label	Test Prediction	Test top 3
5	5	2	2	[2, 1, 8],
9	9	5	5	[5, 6, 8],
1	1	1	1	[1, 2, 3],
6	6	7	7	[7, 2, 3],
7	7	7	7	[7, 8, 9],
2	2	3	3	[3, 5, 8],
4	4	5	5	[5, 8, 3],
1	1	8	5	[5, 8, 9],
9	9	6	6	[6, 5, 8],
0	0	7	7	[7, 9, 3],
5	5	0	0	[0, 8, 9],
2	2	2	2	[2, 7, 1],
2	8	0	0	[0, 6, 8],
5	5	5	9	[9, 5, 8],
6	6	7	7	[7, 3, 2],
5	5	1	1	[1, 2, 3]
1	1			
5	5			
3	3			
3	3			
5	5			
6	6			
8	8			
1	1			
6	6			
5	5			
8	8			
3	3			
2	2			
0	0			
0	0			
6	0			

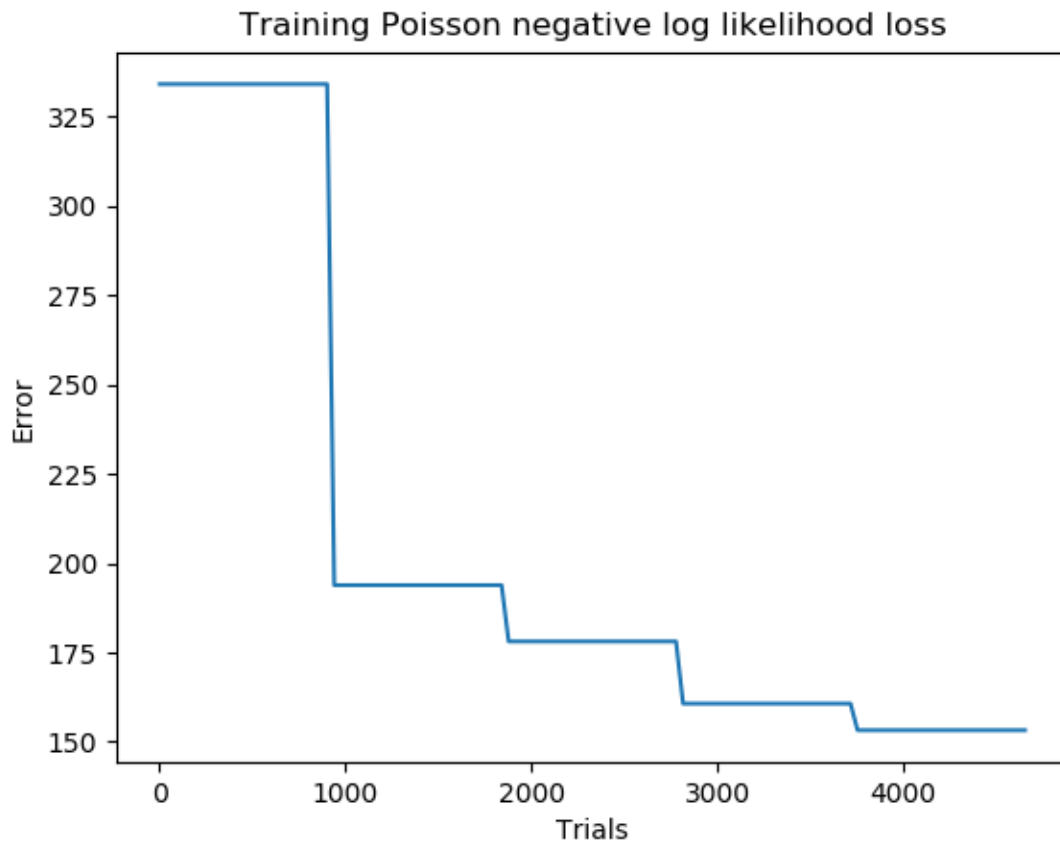
Accuracy Table

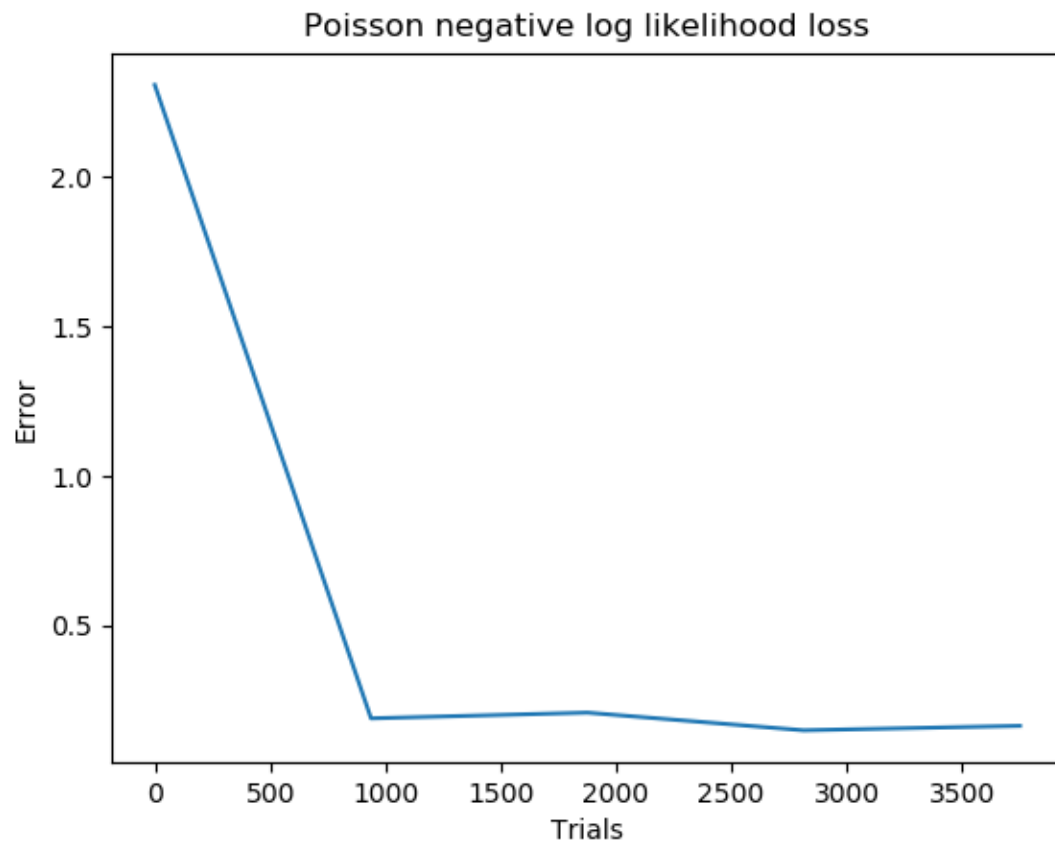
Class	0	1	2	3	4	5	6	7	8	9	Mean
Training	100	100	75	100	100	100	80	100	100	100	95.5
Test	96.32	98.67	95.15	95.34	93.38	96.86	96.34	96.49	94.76	95.83	95.91

Hyper Parameters

Learning Rate= 0.01	Seed: 1	Layering:
Time to Train: 489.388 s	Epochs: 5	Conv, relu, pool, conv, relu, pool, view, linear, relu, linear, softmax
Batch Size: 64	Optimizer: Adam	

Training and Validation loss





Discussion

I started by designing the network by creating the layers; examples called for repeated calls of convolution, relu and pooling. I then reshape the output then linearize in two steps down to the 10 classes I then use a softmax function to classify each class. I then optimize using Adam (AdaDelta with momentum) to avoid local minimums due to similar looking structures. I then perform the forward pass and check the loss using a Poisson negative log likelihood model (this takes into account that the output of the forward pass is a logarithmic softmax).

By using a more intense layering structure I was able to use less epochs, smaller batches, and overall decrease the training time while maintaining a high degree of accuracy. Finally, the training data was split-up to have a 64 batche size of shuffled data and validated against all of the training data with test data downloaded separately.

Problem 2: CIFAR

Confusion Table

Train Label	Train Predicted	Test Label	Test Prediction	Test top 3	Test Classification	Test Prediction Classification
699	699	9	1	[1, 4, 9],	Truck	Car
096	096	7	7	[7, 4, 9],	Horse	Horse
472	772	8	8	[8, 0, 9],	Ship	Ship
317	517	5	4	[4, 5, 3],	Dog	Deer
189	189	7	7	[7, 4, 5],	Horse	Horse
355	755	3	5	[5, 1, 3],	Cat	Dog
045	043	8	8	[8, 0, 9],	Ship	Ship
061	061	1	1	[1, 8, 0],	Car	Car
992	990	2	7	[7, 2, 5],	Bird	Horse
739	239	6	6	[6, 4, 3],	Frog	Frog
225	225	0	0	[0, 2, 4],	Plane	Plane
414	414	8	8	[8, 3, 0],	Ship	Ship
018	418	6	6	[6, 2, 3],	Frog	Frog
021	021	5	5	[5, 3, 2],	Dog	Dog
526	526	3	5	[5, 3, 4],	Cat	Dog
824	824	3	2	[2, 0, 8]	Cat	Bird
737	737					
105	005					
372	372					
607	687					
559	379					
395	314					
822	822					
552	752					
417	417					
595	695					
42	45					

Accuracy Table


Translation	plane	car	bird	cat	deer	dog	frog	horse	ship	truck	
Class	0	1	2	3	4	5	6	7	8	9	Mean
Training	71.42	87.5	83.33	66.66	85.71	53.86	100	87.5	100	90	82.59
Test	64.3	77.3	54.4	42.0	60.8	54.6	71.1	68.4	70.2	69.2	63.23

Hyper Parameters

Learning Rate= 0.1	Seed: 1	Layering:
Time to Train: 2689.7 s	Epochs: 36	Conv, relu, pool, view, linear, relu, linear
Batch Size: 128	Optimizer: Adadelata	

Training and Validation loss





Discussion

Like the MNIST data, I started by designing the network by creating the layers; unlike the MNIST data using multiple layer increased the output (before linearization) above 12,000. This drastically increased the training time. I then reshape the output then linearize in two steps down to the 10 classes. I then optimize using AdaDelta. I then perform the forward pass and check the loss using cross entropy. These parameters were changed from the MNIST as using a SoftMax function to aid in classification caused an extreme drop off in accuracy.

In order to compensate for a smaller layering, I increased the number of epochs and the learning rate. After trial and error, 36 epochs were the goldilocks zone before loss began increasing. I also tried the lowering the learning rate to 0.01, however after 100 epochs, I was unable to achieve the same/better results without significantly more training time being needed exceeding the already lengthy training time of 45 minutes. Additionally, compared to the MNIST CNN, a batch size of 128 was used as smaller batch sizes (64) would fail to train on all input data types.

Finally, the training data was split-up to have 128 batch size of shuffled data and validated against all of the training data with test data downloaded separately.

Key takeaways

Several takeaways I took from this are how the Pytorch works and several of the various functions that can be used in CNN's.

I also came to understand how structuring different CNN's can have a significant affect on how well a CNN will perform. This is particularly targeted towards pictures with strong easily quantifiable information (black numerals on a white background) and complex imagery (dogs vs. cats).